

# Post-Hurricane Structure Damage Assessment Leveraging Aerial Imagery and Convolutional Neural Networks

**Allan Kapoor**  
Springboard School of Data  
December 2021



# The Need for Post-Disaster Damage Assessment

Major cause of damage/economic disruption in the United States:

- Hurricanes: \$21.5 billion per event
- 2015-2020: 10+ >\$1 billion storms per year

Damage assessment critical to response efforts but challenging:

- Lack of access
- Manual review is resource and time intensive



# Problem Statement

In order to support effective natural disaster response and recovery, how can we leverage remote sensing imagery to quickly and efficiently produce accurate damage assessments after extreme storms?

**Stakeholders:** Federal Emergency Management Agency (FEMA), the Department of Defense (National Guard, etc.), state offices of emergency services, and county/municipal governments.

# Data Source

University of Washington Disaster Data Science Lab:

<https://ieee-dataport.org/open-access/detecting-damaged-buildings-post-hurricane-satellite-imagery-based-customized>

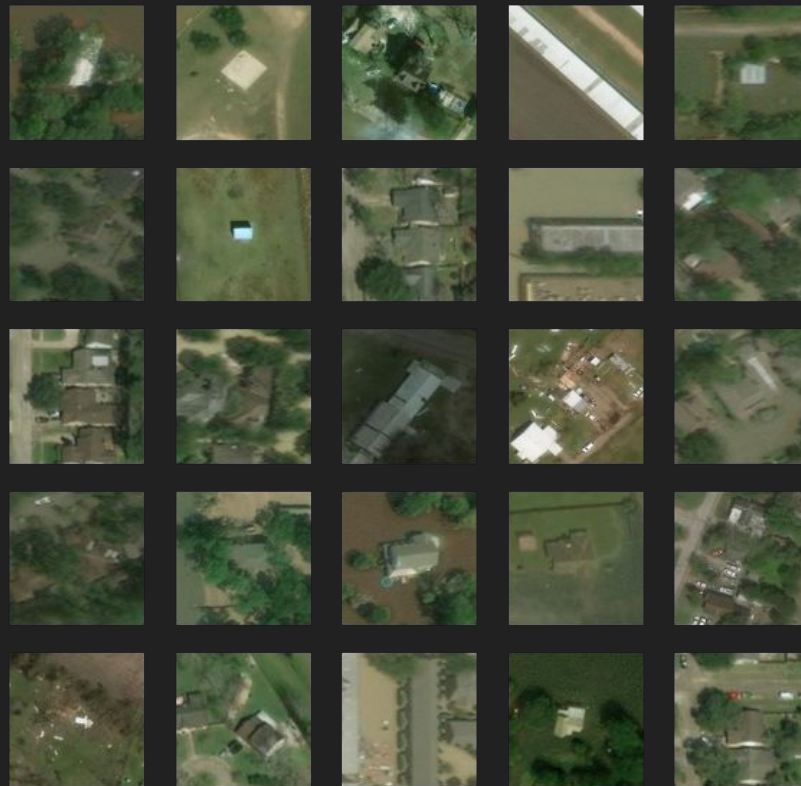
- Aerial images of structures from Houston area after Hurricane Harvey (2017)
- 14,000 images tagged as with damage or no damage
  - Training set: 8,000 images
  - Validation set: 2,000 images
  - Test set: 2,000 images
- Labelling based on community crowdsourced ground-level data collection
- Equal number of images in each class

# Challenges

“Damage” is not a single feature or object

Scattered materials or poor structure condition may be due to dilapidation, not recent damage

Buildings of same class may have very different structures



# Data Exploration: Visual inspection

Color, RGB format

128x128 pixels

Observations:

- Flooding: texture
- Debris
- Difficult task to human eye

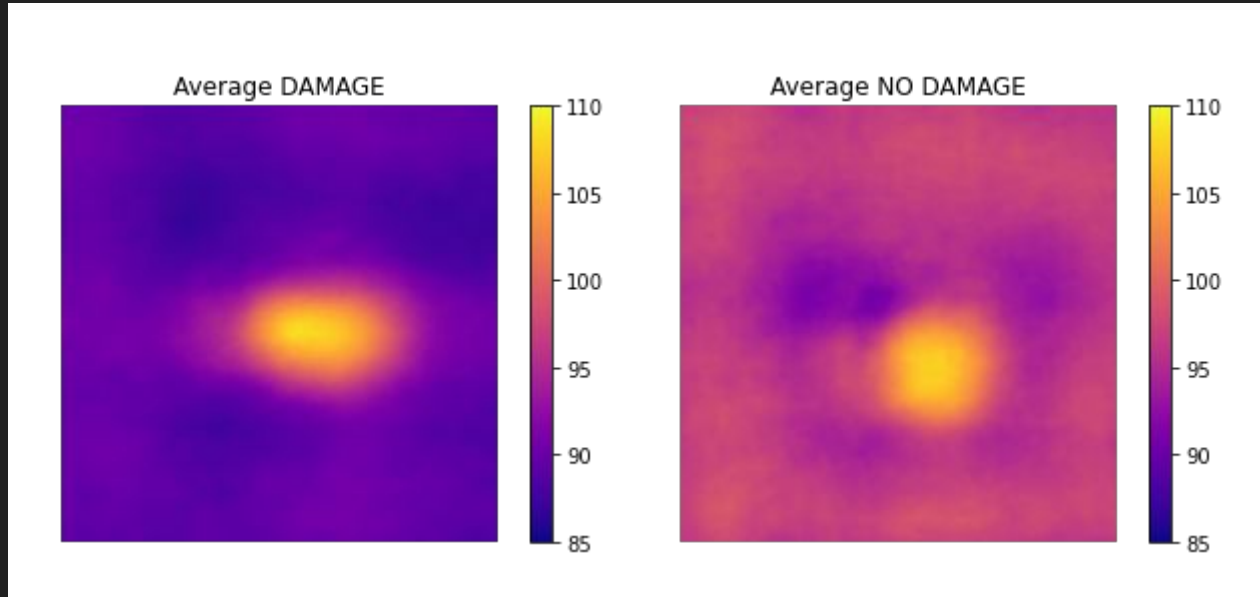


# Data Exploration: Trends by Class

## Mean value for each pixel by class

Structure apparent in both

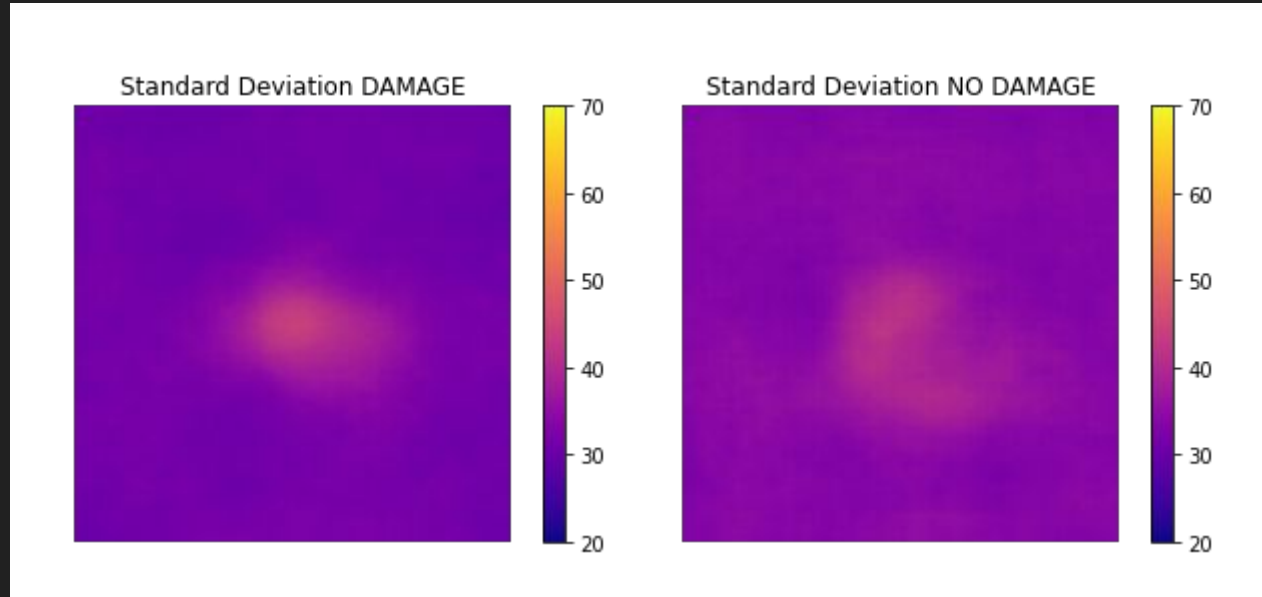
Lower value pixels surrounding structure in average image with damage than without



# Data Exploration: Trends by Class

## Standard deviation for each pixel by class

Variation around structure greater for no damage - because ground is visible above flood water?

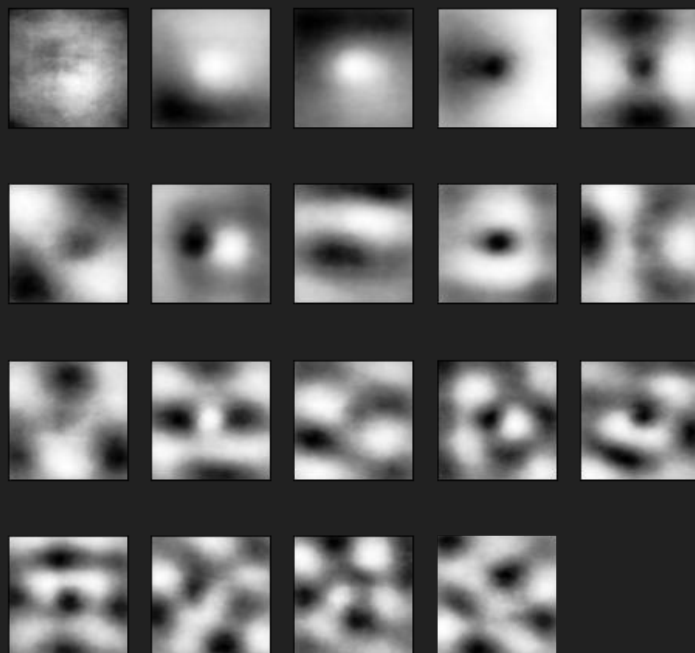




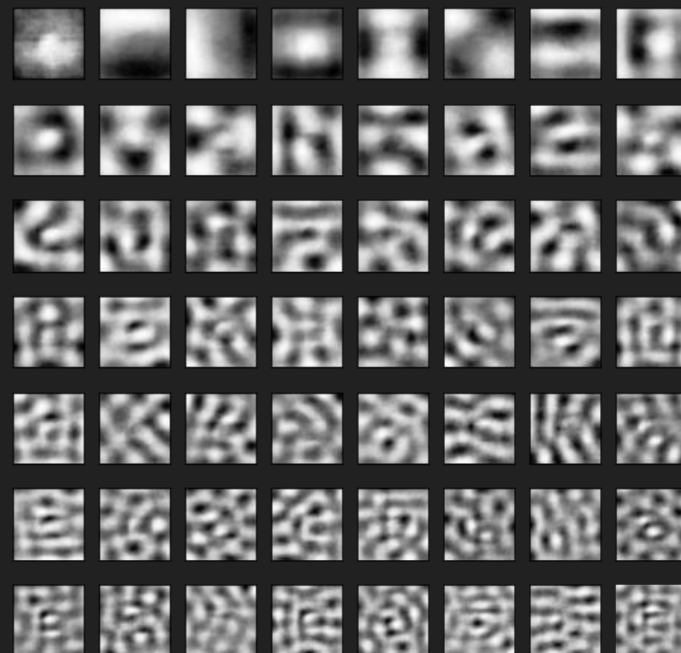
# Data Exploration: PCA/Eigenimages

**PCA**  
**explaining**  
**70% of**  
**variation**

With damage: 19 principle components



With damage: 56 principle components



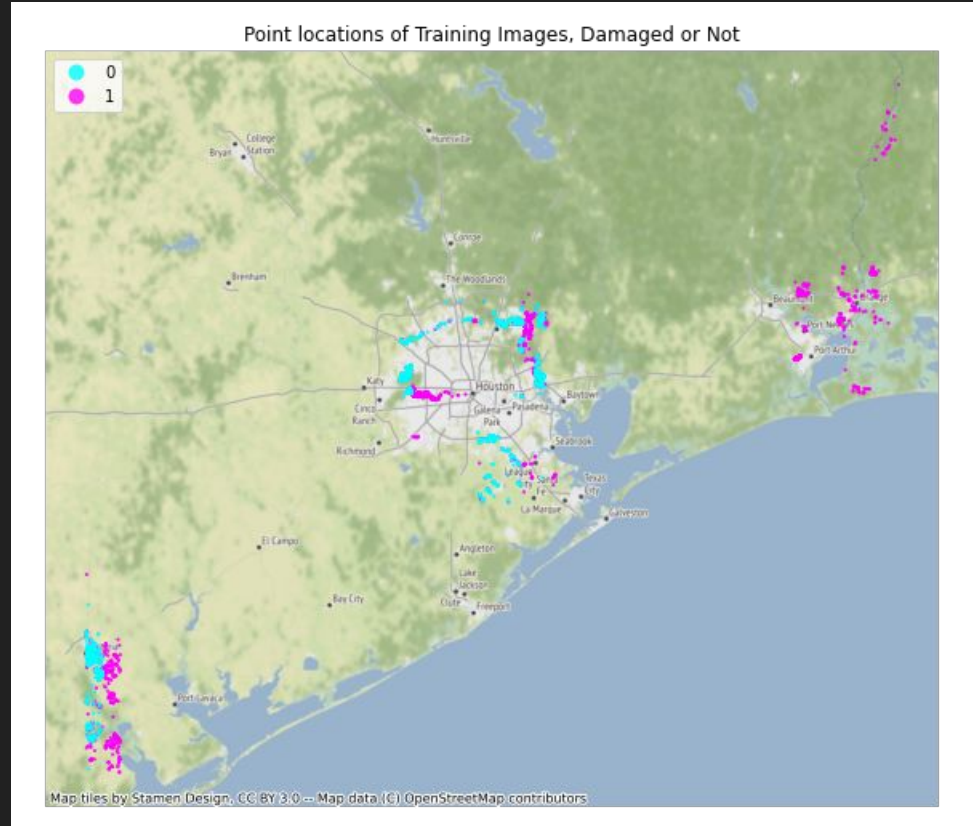
## Data Exploration: Geographic Distribution

## Geographic distribution of training data by class

(0=damage, 1=no damage)

## 3 distinct areas

Spatial concentrations of damaged and not damaged structures - may result in training on irrelevant features?



# Modeling Approach

Convolutional neural network

TensorFlow (Keras API)

Performance metric: accuracy on validation set

Google Cloud VM: 8 vCPUs (30 GB RAM) + 1 NVIDIA T4 GPU (16 GB)

Iterations:

- Baseline Model
- Add max pooling, drop layers
- Transfer learning

# Modeling: Baseline

3 convolution layers, 3 dense layers

- Batch normalization
- ReLU activation
- Adam optimizer

Accuracy (validation set):

- **0.94650** (no image augmentation)
- **0.95650** (with image augmentation)

Layer	Output Shape	# of Params
Rescaling	(128, 128, 3)	0
Convolution (filters=32, kernel_size=5, strides=2)	(64, 64, 32)	2,432
Batch Normalization	(64, 64, 32)	128
Activation (ReLU)	(64, 64, 32)	0
Convolution (filters=32, kernel_size=3, strides=1)	(64, 64, 64)	18,496
Batch Normalization	(64, 64, 64)	256
Activation (ReLU)	(64, 64, 64)	0
Convolution (filters=32, kernel_size=3, strides=1)	(64, 64, 64)	36,928
Batch Normalization	(64, 64, 64)	256
Activation (ReLU)	(64, 64, 64)	0
Flattening	262,144	0
Dense (512 nodes, ReLU activation)	512	134,218,240
Dense (256 nodes, ReLU activation)	256	131,328
Dense (124 nodes, ReLU activation)	124	32,896
Dense (2 nodes, Softmax activation)	2	258

# Modeling: Improving Performance

Reducing overfitting:

Max Pooling layers (reduce spatial sensitivity of convolution filters)

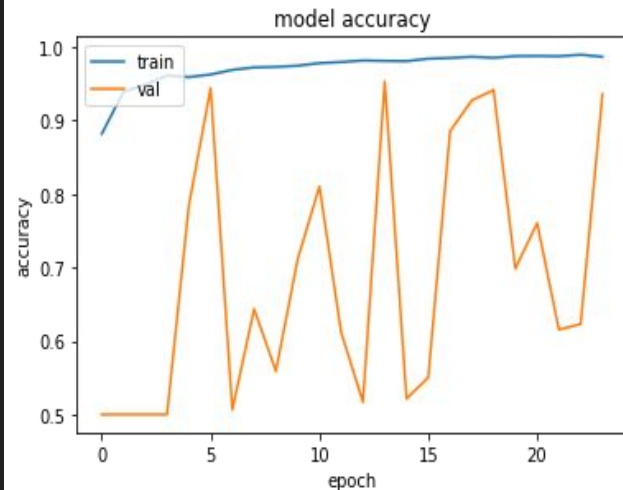
Drop out layers (reduce overfitting to noise)



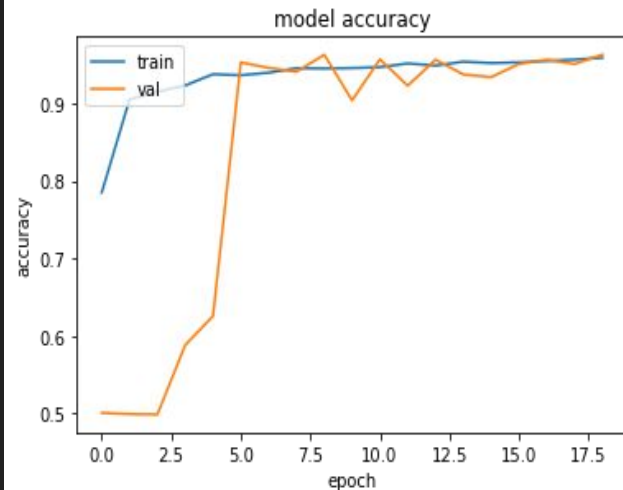
# Modeling: Improving Stability/Convergence

- Reduction in kernel size and stride for the first convolution layer
- Reduction in number of filters in first convolution layer
- Reduction in number of nodes in each dense layer by 50%
- Reduction in initial learning rate for the Adam optimizer from the default (0.001) to 0.0001

**Not Converging:**  
(before updates)



**Converging:**  
(after updates)



# Modeling: Refined Model Architecture

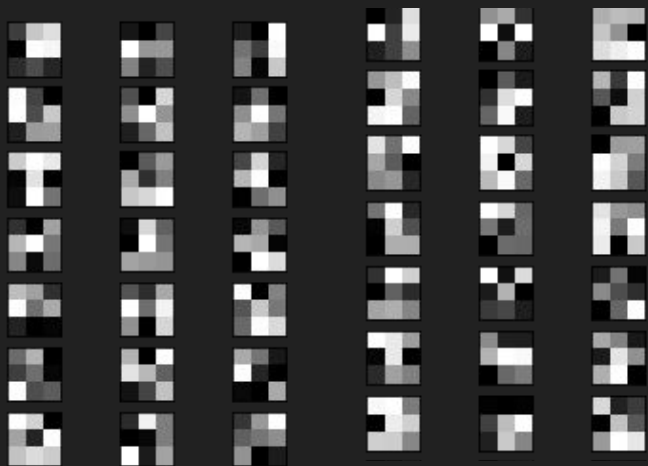
Validation accuracy: **0.9735**

Layer	Output Shape	# of Params
Rescaling	(128, 128, 3)	0
Convolution (filters=32, kernel_size=3, strides=1)	(128, 128, 32)	896
Max Pooling (pool size=2, strides=2)	(64, 64, 32)	0
Batch Normalization	(64, 64, 32)	128
Activation (ReLU)	(64, 64, 32)	0
Convolution (filters=32, kernel_size=3, strides=1)	(32, 32, 64)	18,496
Max Pooling (pool size=2, strides=2)	(16, 16, 64)	0
Batch Normalization	(16, 16, 64)	256
Activation (ReLU)	(16, 16, 64)	0
Convolution (filters=32, kernel_size=3, strides=1)	(8, 8, 64)	36,928
Max Pooling (pool size=2, strides=2)	(4, 4, 64)	0
Batch Normalization	(4, 4, 64)	256
Activation (ReLU)	(4, 4, 64)	0
Flattening	1024	0
Dense (512 nodes, ReLU activation)	512	524,800
Dropout (rate=0.3)	512	0
Dense (256 nodes, ReLU activation)	256	131,328
Dropout (rate=0.2)	256	0
Dense (128 nodes, ReLU activation)	128	32,896
Dropout (rate=0.1)	128	0
Dense (2 nodes, Softmax activation)	2	258

# Modeling: Convolution Filters

Smaller kernel size (3) performed better than medium (5) or large (10)

**Kernel size=3** (best performance)



**Kernel size=10** (more visible shapes, but poor performance)



# Modeling: Deep Network via Transfer Learning

## Resnet50

- Balance of performance/computational efficiency
- Weights trained on ImageNet
- Used for xView2 competition baseline

## Incorporation into model:

- Convolution/max pooling
- Resnet model architecture (frozen weights)
- Dense layers



# Comparing Models

Deep network achieved highest accuracy, but very slow

More basic model only misclassified 6 more images but classifies much more efficiently

Model	Validation Accuracy
Baseline (no image augmentation)	0.9365
Baseline	0.9440
With Max Pooling (kernel=5) & Dropout Layers	0.9500
With Max Pooling (kernel=10) & Dropout Layers	0.9230
<b>With Max Pooling (kernel=3) &amp; Dropout Layers (dense layers with 50% less nodes)</b>	<b>0.9735</b>
Transfer Learning (with Max Pooling kernel=5)	0.9765
Transfer Learning (with Max Pooling kernel=3)	0.9735

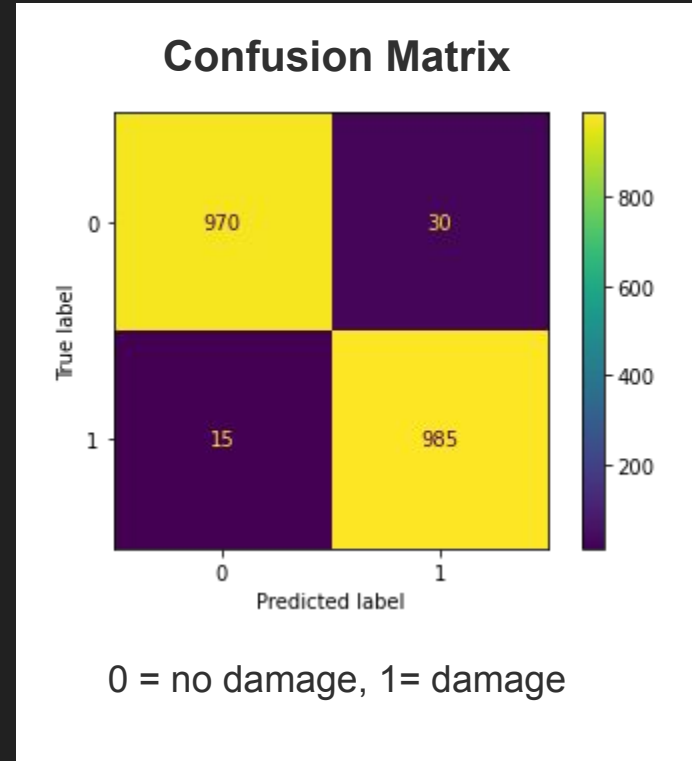


# Selected Model: Performance on Test Data

Accuracy on test data: **0.9775**

**False positive rate:**

**Anything else?**



# Selected Model: Incorrectly Classified Images

## False Positives

(30 total)



## False Negatives

(15 total)

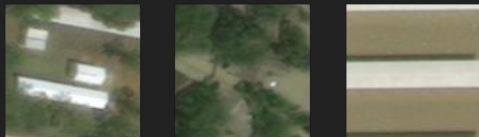


**False Positives** tend to have surfaces that are mistaken for flood waters, or have junk around them that is mistaken for damage. False positives also tend to be rural structures.

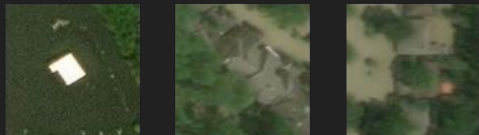
**False negatives** appear to be mostly large or non-residential structures, have a lot of variation in the ground surface, and/or no obvious flood water

## Selected Model: Correctly Classified Images

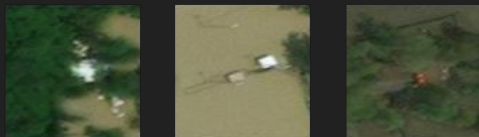
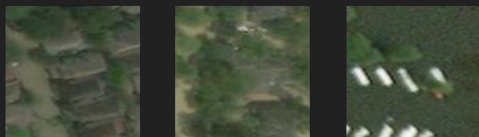
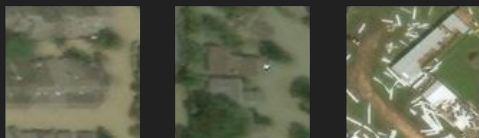
## True Positive Examples



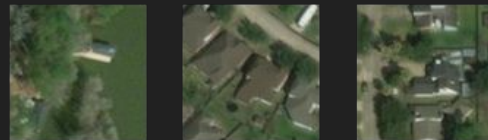
Tend to have obvious  
flood water, or  
scattered materials.



Some of these aren't obvious to the human eye.



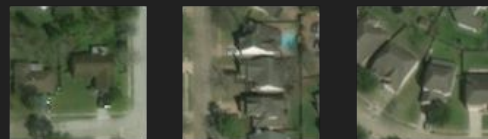
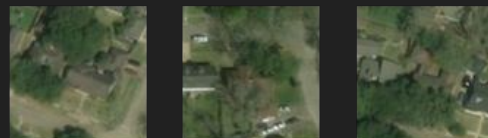
## True Negative Examples



Tend to be obviously  
not flooded (visible  
ground)



Also examples that do have water or surfaces that look like flood water to the human eye.



# Conclusion: Opportunities

Improve training data:

- More labelled data
- Imagery from after floods have subsided
- Imagery from neighborhoods with mixed impacts
- Imagery from different cities

Improve modeling:

- Cross validation
- Hyperparameter optimization

# Conclusion: Using the Model

Damage classification as step in automated pipeline:

- Ingest/clean aerial imagery
- Crop images of structures based on MS Building Footprints data
- Classify images
- Plot locations/damage assessment on interactive map

Much faster than crowdsourcing/manual review: classified 2,0000 structures in seconds on desktop computer



**Hurricane Ida - Crowdsourced Damage Assessment App**  
327 GIS staff required to ID 197,712 damaged structures